UNIVERSITAT DE BARCELONA

Master in Fundamental Principles of Data Science

Dr Rohit Kumar

UNIVERSITAT DE BARCELONA

Today's Objective

- Introduction to NoSQL
- Introduction to MongoDB

UNIVERSITAT DE BARCELONA

NoSQL

UNIVERSITAT DE BARCELONA

What is NoSQL

Scalable

Fast

Flexible and Easy

UNIVERSITAT DE BARCELONA

SQL to NoSQL

NoSQL => "Not only SQL."

	NoSQL	SQL
Model	Non-relational	Relational
	Stores data in JSON documents, key/value pairs, wide column stores, or graphs	Stores data in a table
Data	Offers flexibility as not every record needs to store the same properties	Great for solutions where every record has the same properties
	New properties can be added on the fly	Adding a new property may require altering schemas or backfilling data
	Relationships are often captured by denormalizing data and presenting it in a single record	Relationships are often captured in a using joins to resolve references across tables
	Good for semi-structured data	Good for structured data
Schema	Dynamic or flexible schemas	Strict schema
	Database is schema-agnostic and the schema is dictated by the application. This allows for agility and highly iterative development	Schema must be maintained and kept in sync between application and database
Transactions	ACID transaction support varies per solution	Supports ACID transactions
Consistency	Consistency varies per solution, some solutions have tunable consistency	Strong consistency supported
Scale	Scales well horizontally	Scales well vertically

Source: https://olive4oyl.wordpress.com/2016/08/31/sql-vs-nosql/

UNIVERSITAT DE BARCELONA

Advantages of NoSQL

- Can be used as Primary or Analytic Data Source
- Big Data Capability
- No Single Point of Failure
- Easy Replication
- No Need for Separate Caching Layer
- It provides fast performance and horizontal scalability.
- Can handle structured, semi-structured, and unstructured data with equal effect
- Object-oriented programming which is easy to use and flexible

UNIVERSITAT DE BARCELONA

Advantages of NoSQL

- NoSQL databases don't need a dedicated high-performance server
- Support Key Developer Languages and Platforms
- Simple to implement than using RDBMS
- It can serve as the primary data source for online applications.
- Handles big data which manages data velocity, variety, volume, and complexity
- Excels at distributed database and multi-data center operations
- Eliminates the need for a specific caching layer to store data
- Offers a flexible schema design which can easily be altered without downtime or service disruption

UNIVERSITAT DE BARCELONA

Disadvantages of NoSQL

- No standardization rules
- Limited query capabilities
- RDBMS databases and tools are comparatively mature
- It does not offer any traditional database capabilities, like consistency when multiple transactions are performed simultaneously.
- Doesn't work as well with relational data
- The learning curve is stiff for new developers
- Open source options so not so popular for enterprises.

SQL vs NoSQL

Src: https://scalegrid.io/blog/2019-database-trends-sql-vs-nosql-top-databases-single-vs-multiple-database-use/

UNIVERSITAT DE BARCELONA

NoSQL

Key Value

Example:
Riak, Tokyo Cabinet, Redis server, Memcached, Scalaris

Document-Based

Example:
MongoDB, CouchDB, OrientDB, RavenDB

Column-Based

Example:
BigTable, Cassandra, Hbase, Hypertable

Graph-Based

Example:
Neo4J, InfoGrid, Infinite Graph, Flock DB

Key Value

This key/value type database allow clients to read and write values using a key as follows:

- Get(key), returns the value associated with the provided key.
- Put(key, value), associates a value with the key.
- Multi-get(key1, key2, .., keyN), returns the list of values associated with the list of keys.
- Delete(key), removes the entry for the key from the data store.

Key	Value
“India”	{“B-25, Sector-58, Noida, India – 201301”}
“Romania”	{“IMPS Moara Business Center, Buftea No. 1, Cluj-Napoca, 400606″,City Business Center, Coriolan Brediceanu No. 10, Building B, Timisoara, 300011”}
“US”	{“3975 Fair Ridge Drive. Suite 200 South, Fairfax, VA 22033”}

Use cases

Session store: A session-oriented application such as a web application starts a session when a user logs in and is active until the user logs out or the session times out. During this period, the application stores all session-related data either in the main memory or in a database. Session data may include user profile information, messages, personalized data and themes, recommendations, targeted promotions, and discounts. Each user session has a unique identifier. Session data is never queried by anything other than a primary key, so a fast key-value store is a better fit for session data. In general, key-value databases may provide smaller per-page overhead than relational databases.

Use cases

During the holiday shopping season, an e-commerce website may receive billions of orders in seconds. Key-value databases can handle the scaling of large amounts of data and extremely high volumes of state changes while servicing millions of simultaneous users through distributed processing and storage. Key-value databases also have built-in redundancy, which can handle the loss of storage nodes.

Document-Based

A document database is a type of nonrelational database that is designed to store and query data as JSON-like documents.

Document databases enable flexible indexing, powerful ad hoc queries, and analytics over collections of documents.

{country:”India”, {Street: “B-25, City:”Noida”, State:”UP”, Pincode:”201301”} }

{country :”Romania”, {Boulevard:”Coriolan Brediceanu No. 10”, Block:”B, Ist Floor”, City: “Timisoara”, Pincode: 300011”} }

{country :”US”, {Latitude:”40.748328”, Longitude:”-73.985560”} }

UNIVERSITAT DE BARCELONA

Use cases

A document database is a great choice for content management applications such as blogs and video platforms. With a document database, each entity that the application tracks can be stored as a single document. The document database is more intuitive for a developer to update an application as the requirements evolve. In addition, if the data model needs to change, only the affected documents need to be updated. No schema update is required and no database downtime is necessary to make the changes.

UNIVERSITAT DE BARCELONA

Use cases

Document databases are efficient and effective for storing catalog information. For example, in an e-commerce application, different products usually have different numbers of attributes. Managing thousands of attributes in relational databases is inefficient, and the reading performance is affected. Using a document database, each product’s attributes can be described in a single document for easy management and faster reading speed. Changing the attributes of one product won’t affect others.

Search Based Engines (kind of document store)

Search Engine Database

- A search-engine database is a type of nonrelational database that is dedicated to the search of data content.
- Search-engine databases use indexes to categorize the similar characteristics among data and facilitate search capability.
- Example: Elastic Search or Splunk

Use Cases

Text search

Search-engine databases can handle full-text search faster than relational databases. For example, an e-commerce website can use search-engine databases to provide instant autocompletes or suggestions for its customers. Search-engine databases can sort relevant results based on characteristics such as name, price, category, or release data, and display the results in a structured view.

Use Cases

Logging and analysis

Maintaining larger applications that are either distributed across several nodes or consist of several smaller applications searching for events in log files can become tedious. Search-engine databases can handle the logging more efficiently. You can centralize your logs from different applications by indexing them using a search-engine database. For example you can see the logs of your Apache web server combined with the log files of your application server. Because all the information is available in real time, you can implement a visual representation of what is happening in your system in real time, which can help you to find problems more quickly.

UNIVERSITAT DE BARCELONA

Column-Based

- A column store database is a type of database that stores data using a column oriented model.
- Columns store databases use a concept called a keyspace. A keyspace is kind of like a schema in the relational model. The keyspace contains all the column families (kind of like tables in the relational model), which contain rows, which contain columns.

UNIVERSITAT DE BARCELONA

Column-Based

Column-Based

- A **column family** consists of multiple rows.
- Each **row** can contain a different number of columns to the other rows. And the columns don't have to match the columns in the other rows (i.e. they can have different column names, data types, etc).
- Each **column** is contained to its row. It doesn't span all rows like in a relational database. Each column contains a name/value pair, along with a timestamp. Note that this example uses Unix/Epoch time for the timestamp.

Column-Based

- **Row Key**. Each row has a unique key, which is a unique identifier for that row.
- **Column**. Each column contains a name, a value, and timestamp.
- **Name**. This is the name of the name/value pair.
- **Value**. This is the value of the name/value pair.
- **Timestamp**. This provides the date and time that the data was inserted. This can be used to determine the most recent version of data.

Graph Based

- Graph databases are purpose-built to store and navigate relationships.
- These databases uses edges and nodes to represent and store data.
- Nodes are organised by some relationships with one another, which is represented by edges between the nodes.
- Both the nodes and the relationships have some defined properties.

Graph Based

Sam
Friend
Annie
Friend
Mac
Friend
Jack
Friend
Doug
Friend
Friend
Harry
Howard

Id: 2
Name: Bob
Age: 22
Id: 100
Label: knows
Since: 2001/10/03
Id: 101
Label: knows
Since: 2001/10/04
Id: 105
Label: is_member
Since: 2011/02/14
Id: 104
Label: Members
Id: 1
Name: Alice
Age: 18
Id: 103
Label: Members
Id: 102
Label: is_member
Since: 2005/7/01
Id: 3
Type: Group
Name: Chess

Use cases

Graph databases are capable of sophisticated **fraud prevention**. With graph databases, you can use relationships to process financial and purchase transactions in near-real time. With fast graph queries, you are able to detect that, for example, a potential purchaser is using the same email address and credit card as included in a known fraud case. Graph databases can also help you easily detect relationship patterns such as multiple people associated with a personal email address, or multiple people sharing the same IP address but residing in different physical addresses.

Use cases

Graph databases are a good choice for recommendation applications. With graph databases, you can store in a graph relationships between information categories such as customer interests, friends, and purchase history. You can use a highly available graph database to make product recommendations to a user based on which products are purchased by others who follow the same sport and have similar purchase history. Or, you can identify people who have a friend in common but don’t yet know each other, and then make a friendship recommendation.

UNIVERSITAT DE BARCELONA

Types of NoSQL

Types	Performance	Scalability	Flexibility	Complexity
Key-Value	high	high	high	None
Column Store	high	high	Moderate	Low
Document Store	high	Variable (high)	high	Low
Graph Database	Variable	Variable	high	high

UNIVERSITAT DE BARCELONA

MongoDB

UNIVERSITAT DE BARCELONA

Document Model

Lets take an example of a post/blog on a social media platform like redit or twitter.

Every post contains following elements:

1) Title of the post
2) Tags assosicated with the post
3) Comments made
4) Votes
5) Author
6) URL of the post
7) Images or video in the post

UNIVERSITAT DE BARCELONA

In Relational Model

In Document Model

```
{
  _id:
ObjectID('4bd9e8e17cefd644108961bb'),
  title: String,
  url: String,
  author: String,
  vote_count: Integer,
  tags: Array[
    String
  ],
  image: ImageObject,
  comments: Array[
    CommentObject
  ]
}
```

ImageObject:

```
{
url: String,
caption: String,
type: String,
size: Integer,
data: String
}
```

CommentObject:

```
{
user:String,
text:String
}
```

In Document Model

```
{
  _id: ObjectID('4bd9e8e17cefd644108961bb'),
  title: 'Adventures in Databases',
  url: 'http://example.com/databases.txt',
  author: 'msmith',
  vote_count: 20,
  tags: [
    'databases',
    'mongodb',
    'indexing'
  ],
  image: {
    url: 'http://example.com/db.jpg',
    caption: 'A database.',
    type: 'jpg',
    size: 75381,
    data: 'Binary'
  },
  comments: [
    {
      user: 'bjones',
      text: 'Interesting article.'
    },
    {
      user: 'sverch',
      text: 'Color me skeptical!'
    }
  ]
}
```

Advantage

Suppose you want to find all posts tagged with the term politics having more than 10 votes.

```
SELECT * FROM posts
INNER JOIN posts_tags ON posts.id = posts_tags.post_id
INNER JOIN tags ON posts_tags.tag_id == tags.id
WHERE tags.text = 'politics' AND posts.vote_count > 10;
```

In MongoDB

```
db.posts.find({'tags': 'politics', 'vote_count': {'$gt': 10}});
```

UNIVERSITAT DE BARCELONA

SQL to Mongo

SQL Terms/Concepts	MongoDB Terms/Concepts
database	database
table	collection
row	document
column	field
index	index
table joins	$lookup, embedded documents
primary key Specify any unique column or column combination as primary key.	primary key In MongoDB, the primary key is automatically set to the __id field.

UNIVERSITAT DE BARCELONA

SQL to Mongo

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	$match
GROUP BY	$group
HAVING	$match
SELECT	$project
ORDER BY	$sort
LIMIT	$limit
SUM()	$sum
COUNT()	$sum $sortByCount
join	$lookup
SELECT INTO NEW_TABLE	$out

UNIVERSITAT DE BARCELONA

High Availability in Mongo

MongoDB's replication facility, called replica set, provides:

- automatic failover
- data redundancy.

A replica set is a group of MongoDB servers that maintain the same data set, providing redundancy and increasing data availability.

UNIVERSITAT DE BARCELONA

Replication

Replication

Automatic Failover

UNIVERSITAT DE BARCELONA

Why Replication?

- To keep your data safe
- High (24*7) availability of data
- Disaster recovery
- No downtime for maintenance (like backups, index rebuilds, compaction)
- Read scaling (extra copies to read from)
- Replica set is transparent to the application

Sharding

Sharding is a method for distributing data across multiple machines. MongoDB uses sharding to support deployments with very large data sets and high throughput operations.

Sharding

A MongoDB sharded cluster consists of the following components:

- **shard**: Each shard contains a subset of the sharded data. Each shard can be deployed as a replica set.
- **mongos**: The mongos acts as a query router, providing an interface between client applications and the sharded cluster.
- **config servers**: Config servers store metadata and configuration settings for the cluster. As of MongoDB 3.4, config servers must be deployed as a replica set (CSRS).

Sharding

App Server
Router
(mongos)
App Server
Router
(mongos)
1 or more Routers
Config Servers
(replica set)
2 or more Shards
Shard
(replica set)
Shard
(replica set)

UNIVERSITAT DE BARCELONA

Sharding

Sharding Strategy

Hashed Sharding involves computing a hash of the shard key field's value. Each chunk is then assigned a range based on the hashed shard key values.

Sharding

Ranged sharding involves dividing data into ranges based on the shard key values. Each chunk is then assigned a range based on the shard key values.

Sharded Cluster Balancer

UNIVERSITAT DE BARCELONA

Before Next week

1) Run the following command and start mongoDB using docker

docker run -p 27017:27017 -v <an empty folder path to store data>:/data/db mongo

Make sure to replace ***<an empty folder path to store data>*** with an actual folder path in your disk.

2) Install MongoDB Compass

https://docs.mongodb.com/compass/master/install/